

1 Regular Expressions and Finite State Automata

- 1.1 Finite State Automaton**
- Formal definition: $(Q, \Sigma, q_0, F, \delta(q, i))$ where:
 - Q : Set of states
 - σ : input alphabet
 - q_0 : the start state
 - F : the set of final state
 - $\delta(q, i)$: transition function between the states
 - Deterministic Finite State Automaton**
 - Accepts an input string if we run out of input and it is in an accepting state
 - Exactly one transition for a given symbol from one state to another
 - Non-Deterministic Finite State Automaton**
 - Accepts an input string if there is at least some path to an accepting state that exhausts the input string
 - Could have more than one transition for the same input symbol or transition to another empty state through an empty symbol
 - Convert from NFSA to DFSA**
 - Keep track of all states that can be transitioned from NFSA
 - A state in converted DFSA denotes a combination of states in NFSA

2 Words, Spelling Errors and Edit Distance

- 2.1 Porter Stemming Algorithm**
- ATONAL -> ATE
 - ING -> ϵ
 - SSES -> SS
- 2.2 Bayesian Classification**
- $$c \hat{=} \operatorname{argmax}_{c \in C} P(c|o) = \operatorname{argmax}_{c \in C} \frac{P(o|c).P(c)}{P(o)} = \operatorname{argmax}_{c \in C} P(o|c).P(c)$$

where $P(o|c)$ is the likelihood and $P(c)$ is the prior

2.3 Minimum Edit Distance

function MIN-EDIT-DISTANCE(source, target) returns min-distance

$m \leftarrow \text{LENGTH}(\text{source})$

$n \leftarrow \text{LENGTH}(\text{target})$

Create a matrix distance[m + 1, n + 1]

distance[0, 0] \leftarrow 0

for each row i from 1 to m do

 distance[i, 0] \leftarrow distance[i - 1, 0] + del-cost(source_i)

for each column j from 1 to n do

 distance[0, j] \leftarrow distance[0, j - 1] + ins-cost(target_j)

for each row i from 1 to m do

 for each column j from 1 to n do

 distance[i, j] \leftarrow MIN(distance[i - 1, j] + del-cost(source_i),

 distance[i - 1, j - 1] + sub-cost(source_i, target_j),

 distance[i, j - 1] + ins-cost(target_j))

return distance[m, n]

3 N-grams

- 3.1 N-gram approximation**
- $P(w_k|w_1, \dots, w_{k-1}) \approx P(w_k|w_{k-(n-1)}, \dots, w_{k-2}, w_{k-1})$
 - MLE: $P(w_k|w_{k-(n-1)}, \dots, w_{k-2}, w_{k-1}, w_k) = \frac{C(w_{k-(n-1)}, \dots, w_{k-2}, w_{k-1}, w_k)}{C(w_{k-(n-1)}, \dots, w_{k-2}, w_{k-1})}$
 - bigram: n = 2 and trigram: n = 3
- 3.2 Perplexity**
- $PP = m(w_1, \dots, w_n)^{-\frac{1}{n}} = \frac{1}{\sqrt[n]{m(w_1, \dots, w_n)}}$
 - For bigrams, $m(w_1, \dots, w_n) = \prod_{k=1}^n nP(w_k|w_{k-1})$
 - Weighted average number of choices a random variable has to make (weighted average branching factor of a language)
- 3.3 Smoothing**
- Add k smoothing: $P(w|w_0) = \frac{C(w_0w) + k}{C(w_0) + kV}$
 - Discount: $\frac{C(w_0w)}{C(w_0w)}$
 - Witten-Bell Smoothing: $P(w|w_0) = \begin{cases} \frac{C(w_0w)}{C(w_0) + T(w_0)} & \text{if } C(w_0w) > 0 \\ \frac{T(w_0)}{Z(w_0)(C(w_0) + T(w_0))} & \text{if } C(w_0w) = 0 \end{cases}$
 - Kneser-Ney Smoothing for Bigram: $P(w|w_0) = \begin{cases} \frac{C(w_0w) - D}{C(w_0)} & \text{if } C(w_0w) > 0 \\ \frac{\alpha(w).[w':C(w'w_0) > 0]}{\sum_w [w':C(w'w_0) > 0]} & \text{if } C(w_0w) = 0 \end{cases}$
- 3.4 Entropy**
- $H(x) = -\sum_{x \in X} p(x) \log_2 p(x)$
 - Entropy for a language: $H(L) = \lim_{n \rightarrow \infty} \inf \frac{1}{n} \log_2 p(w_1, \dots, w_n)$
 - Cross Entropy: $H(p, m) = \lim_{n \rightarrow \infty} \inf \frac{1}{n} \sum_{w \in L} p(w_1, \dots, w_n) \log_2 m(w_1, \dots, w_n)$

- Relation between Perplexity and Entropy:
 $Perplexity = 2^H = 2^{-\frac{1}{n} \log_2 m(w_1, \dots, w_n)} = m(w_1, \dots, w_n)^{\frac{1}{n}}$
- 4 POS Tagging**
 - 4.1 Stochastic POS Tagging**
 - $T \hat{=} \operatorname{argmax}_{t_1, \dots, t_T} \prod_{i=1}^T P(t_i|t_{i-1}).P(w_i|t_i).P(</s>|t_T)$
 - Time Complexity: $O(T.N^T)$
 - 4.2 Viterbi (Dynamic Programming)**
 - $v_t(j) = \max_{i \geq 1} N a_{ij} b_j(o(t))$ where a_{ij} is the transition probability, b_j is the observation likelihood of observation $o(t)$
 - Time Complexity: $O(N^2)$
 - 4.3 Penn Treebank POS Tags**

Tag	Description	Example	Tag	Description	Example
CC	coordin. conjunction	and, but, or	SYM	symbol	%, &
CD	cardinal number	one, two, three	TO	"to"	to
DT	determiner	a, the	UH	interjection	ah, oops
EX	existential 'there'	there	VB	verb, base form	eat
FW	foreign word	mea culpa	VBD	verb, past tense	ate
IN	preposition/sub-conj	of, in, by	VBG	verb, gerund	eating
JJ	adjective	yellow	VBN	verb, past participle	eaten
JJR	adj., comparative	bigger	VBP	verb, non-3sg pres	eat
JJS	adj., superlative	wildest	VBZ	verb, 3sg pres	eats
LS	list item marker	1, 2, One	WDT	wh-determiner	which, that
MD	modal	can, should	WP	wh-pronoun	what, who
NN	noun, sing. or mass	llama	WPS	possessive wh-	whose
NNS	noun, plural	llamas	WRB	wh-adverb	how, where
NNP	proper noun, singular	IBM	\$	dollar sign	\$
NNPS	proper noun, plural	Carolinas	#	pound sign	#
PDT	predeterminer	all, both	"	left quote	" or "
POS	possessive ending	's	"	right quote	" or "
PRP	personal pronoun	I, you, he	(left parenthesis	[, (, {, <
PRPS	possessive pronoun	your, one's)	right parenthesis], }, >
RB	adverb	quickly, never	,	comma	,
RBR	adverb, comparative	faster	.	sentence-final punc	! ;
RBS	adverb, superlative	fastest	:	mid-sentence punc	: ; ... --
RP	particle	up, off	:		


5 Neural Networks

- 5.1 Multi Layer Perceptron**
- $$N_{MLP2}(x) = (g^2(g^1(xW^1 + b^1)W^2 + b^2))W^3 + b^3$$
- $$h^1 = g^1(xW^1 + b^1)$$
- $$h^2 = g^2(xW^2 + b^2)$$
- $$y = h^2W^3 + b^3$$
- 5.2 Activation Functions**
- RELU(x) = $g(x) = \max(0, x)$
 - Sigmoid: $\sigma(x) = \frac{1}{1+e^{-x}}, \sigma'(x) = \sigma(x)(1 - \sigma(x))$
 - Tanh: $\tanh(x) = \frac{e^x + e^{-x}}{e^x - e^{-x}}, \tanh'(x) = 1 - (\tanh(x))^2$
 - HardTanh: $\operatorname{hardtanh}(x) = \begin{cases} -1 & x < -1 \\ x & x > 1 \\ 1 & \text{otherwise} \end{cases}$
- 5.3 Loss Functions**
- Cross Entropy Loss: $L_{\text{logistic}}(\hat{y}, y) = -y \log_2(\hat{y}) - (1 - y) \log_2(1 - \hat{y})$ where $y \in \{0, 1\}$ for binary classification
 - For multi class classification ($C > 2$), $L_{\text{logistic}}(\hat{y}, y) = -\sum_j C_j t_j \log_2(\hat{y}_t)$
 - Softmax loss: $L_{\text{softmax}}(\hat{y}, y) = \frac{e^{\hat{y}[i]}}{\sum_j e^{\hat{y}[j]}}$, squishes range of values to (0, 1), a probability distribution
 - Squared (quadratic) loss: $L_{\text{squared}}(y/hat, y) = 0.5 * (y - y)^2$

- 5.4 Gradient Descent**
- $$w_i = w_i - \alpha \frac{\delta L}{w_i} \frac{\delta L}{w_m} = \sum_{i=1}^N N \frac{\delta L}{\delta x_i} \frac{\delta x_i}{\delta w_m}$$
- 6 Word Embeddings**
- 6.1 Word2Vec**
- $s(w, c) = w.c$
- $$P(+|t, c) = \frac{1}{1 + e^{-t.c}}$$
- $$P(-|t, c) = 1 - P(+|t, c) = \frac{e^{-t.c}}{1 + e^{-t.c}}$$


- 6.2 Choosing noise words**
- $P_A(w) = \frac{\text{count}(w)^{\alpha}}{\sum_w \text{count}(w)^{\alpha}}$ where count is the unigram frequency. In practice, \$alpha\$ value of 0.75 works quite well.
 - Maximize objective function:
 $\sum_{(w,c) \in D} \log_2 P(+|t, c) + \sum_{(w,c) \in D} \log_2 P(-|t, c)$

- 6.3 CBOW**
- $c = \sum_i c_i$
 - $\log P(+|w, c_{1:k}) = \log \frac{1}{1 + e^{(\sum_{i=1..k} w.c_i)}}$
- 6.4 Skip-gram**
- $P(+|w, c_{1..k}) = \prod_{i=1}^k \frac{1}{1 + e^{-w.c_i}}$
 - $P(+|w, c_{1..k}) = \sum_{i=1}^k \log \frac{1}{1 + e^{-w.c_i}}$
- 6.5 Cosine similarity**
- $\operatorname{sim}_{\cos}(u, v) = \frac{u.v}{||u||_2 ||v||_2}$
 - $\operatorname{analogy}(m : w - > k : ?) = \operatorname{argmax}_{v \in V \setminus \{m, w, k\}} \cos(v, k - m + w)$


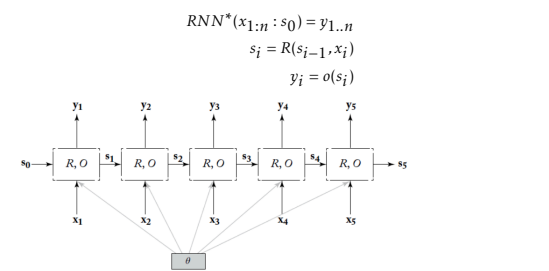


7 CNN

- 7.1 Why is it used?**
- Identify informative ngrams
 - Consider local ordering patterns
- 7.2 Process**
- Apply non linear learned function (filter) over each k-word sliding window
 - Apply l filters to get l-dimensional vector
 - Combine vectors from different windows using pooling into single l-dimensional vector
 - Feed single l-dimensional vector into neural network for prediction
- 7.3 Convolution**
- Narrow convolution (no padding): m = n - (k - 1)
 - Wide convolution (padding k - 1 to each side): m = n + (k - 1)
- 7.4 Pooling**
- Max pooling: $c_j = \max_{1 \leq i \leq m} p_{i[j]} \forall j \in [1, l]$
 - Average pooling: $\frac{1}{m} \sum_{i=1}^m p_i$




- Narrow convolution, n = 9, d = 2, k = 3, l = 3, m = 7
- 8 RNN**
- 8.1 Why is it used?**
- Capture subtle patterns and regularities in sequences
 - Model non-Markovian dependencies
- 8.2 RNN Abstraction**



- 8.3 Elman RNN**
- Sensitive to the order of the words
- $$y_i = O_{RNN}(s_i)$$
- $$x_i \in R^{dx}, s_i \in R^{ds}, W \in R^{(dx+ds).ds}, b \in R_s^{ds}$$
- 8.4 CBOW RNN**
- Does not take into account order of words
- $$s_i = R_{CBOW}(s_{i-1}, x_i) = s_{i-1} + x_i$$
- $$y_i = O_{CBOW}(s_i)$$

9 Seq2Seq



- $$s_i = R_{dec}(s_{j-1}, [t_j^i; c_j^i])$$
- $$y_j = O_{dev}(s_j)$$
- $$P(t_{j+1} | \hat{1}_{1:j}, x_{1:n}) = \operatorname{softmax}(\operatorname{MLP}^{out}(y_j))$$
- $$c_{1:n} = \operatorname{biRNN}_{enc}^*(x_{1:n})$$
- $$\alpha_{[i]}^j = v.tanh([s_{j-1}; c_i]U + b)$$
- $$\alpha_{[i]}^j = \operatorname{softmax}(\alpha_{[1]}^j, \dots, \alpha_{[n]}^j)$$
- $$c^j = \sigma_{i=1}^n \alpha_{[i]}^j . c_i(\operatorname{attend})$$

10 Grammars

- 10.1 CFG**
- $G = (N, \Sigma, P, S)$
- N - terminal symbols
 - Σ - non-terminal symbols
 - P - a -> A where $a \in N, A \in (\Sigma \cup N)^*$
- 10.2 CNF**
- A -> BC or A -> a
 - No epsilon
- 10.3 Equivalence**
- Strong equivalence: $L(G1) = L(G2)$ and same phrase structure
 - Weak equivalence: $L(G1) = L(G2)$ but different phrase structures
- 10.4 Convert to CNF**
- Copy all conforming rules to the new grammar unchanged
 - Convert terminals within rules to dummy non-terminals
 - Convert unit-productions
 - Binarize all rules and add to new grammar
- 10.5 Converting Phrase Structure to Untyped Dependency**
- Find the head (head child is underlined> and pass it up the tree
- 10.6 Parsing**
- Top down parsing: Goal directed, begin from start symbol and then derive parse tree for the given sentence consisting of terminals

5